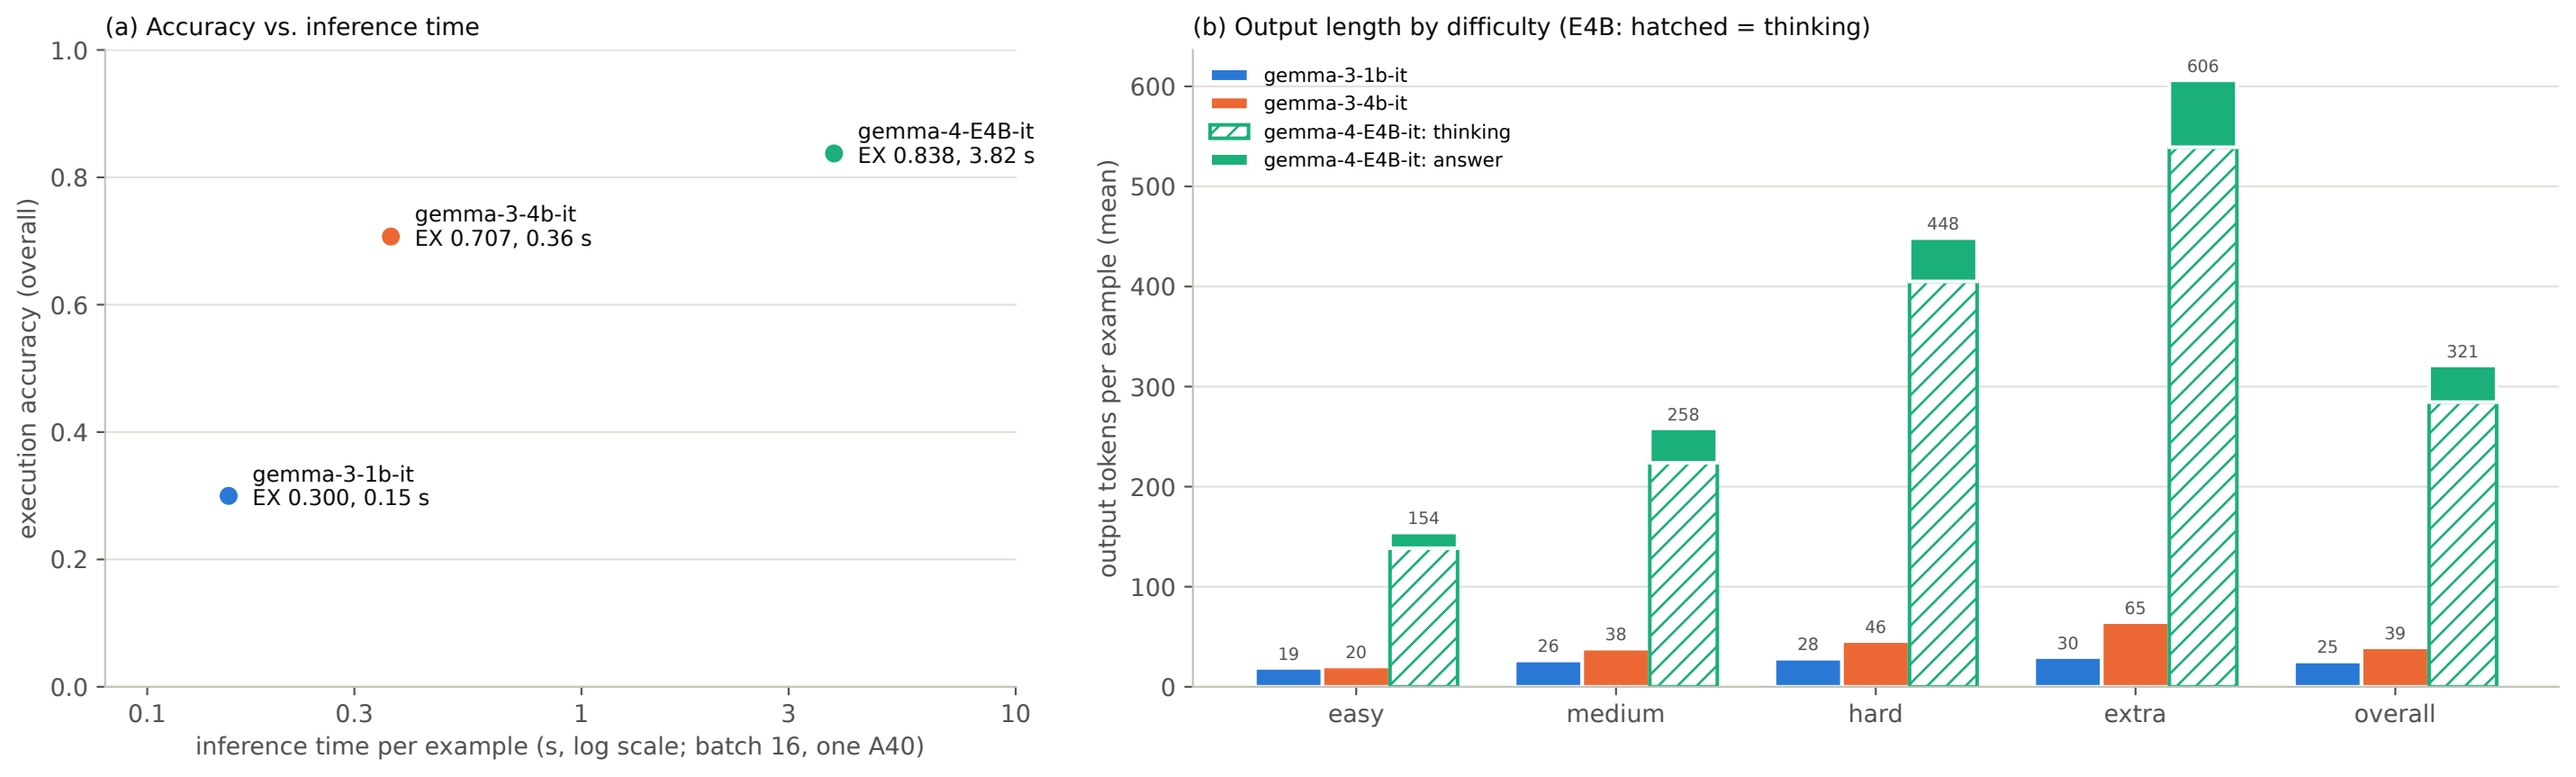
(a) Accuracy vs. inference time
execution accuracy (overall)
inference time per example (s, log scale; batch 16, one A40)
gemma-4-E4B-it
EX 0.838, 3.82 s
gemma-3-4b-it
EX 0.707, 0.36 s
gemma-3-1b-it
EX 0.300, 0.15 s
(b) Output length by difficulty (E4B: hatched = thinking)
output tokens per example (mean)
gemma-3-1b-it
gemma-3-4b-it
gemma-4-E4B-it: thinking
gemma-4-E4B-it: answer
easy
medium
hard
extra
overall